

# 1 Predicting the Airbnb Lodgings in New York

## 1.1 Summary

Using a dataset provided by Airbnb, analysis and predictions will be made to understand what effects the total price of an Airbnb. Diving further, questions like, where do people tend to stay most?, how long do they stay there?, does booking at a certain time of year effect price?, where would be the most affordable place to book a Airbnb?,and what is the average price for a stay? will be addressed. Regressor tree models will be used to predict price of future Airbnbs as well. In the end some results were to be expected but some also came with a surprise. It was no surprise that people stayed the longest in Manhattan or that service fees affected the total price the most. But what was surprising to me at least was that Bedford-Stuyvesant was the neighborhood that people stayed at the most. Upper East Side, Astoria, Prospect-Lefferts Gardens, and East Village were the locations that deemed the most affordable. Finally, people like to stay on average 1-2 nights at Airbnb.

## 2 Import necessary libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
plt.style.use('seaborn')
%matplotlib inline
import seaborn as sns
sns.set_style('darkgrid')
import math
from datetime import datetime
from sklearn import svm
from sklearn import tree

from sklearn.linear_model import LinearRegression
from statsmodels.formula.api import ols
import statsmodels.formula.api as smf
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.model_selection import cross_val_score
from scipy.stats import zscore
import statsmodels.api as sm
import scipy.stats as stats
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.preprocessing import StandardScaler, MinMaxScaler, OneHotEncoder, LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import auc
from sklearn.metrics import accuracy_score, roc_curve, classification_report
from sklearn.metrics import confusion_matrix, plot_confusion_matrix
from sklearn.metrics import recall_score
from sklearn.metrics import precision_score
from sklearn.metrics import f1_score
from sklearn.metrics import roc_auc_score
from sklearn.metrics import mean_squared_log_error
from imblearn.over_sampling import SMOTE, ADASYN
from sklearn.ensemble import RandomForestRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn import metrics
import warnings
warnings.filterwarnings('ignore')
```

Loading the dataset I got from Kaggle in a variable called open\_data.

```
In [2]: open_data = pd.read_csv('data/Airbnb_Open_Data.csv')
```

▼

## 2.1 Exploring Dataset

Looking into the dataset to get a preview of what I will be working with.

```
In [3]: open_data.head()
```

Out[3]:

	id	NAME	host id	host_identity_verified	host name	neighbourhood group	neighbourhood	lat	lon
0	1001254	Clean & quiet apt home by the park	80014485718	unconfirmed	Madaline	Brooklyn	Kensington	40.64749	-73.97
1	1002102	Skylit Midtown Castle	52335172823	verified	Jenna	Manhattan	Midtown	40.75362	-73.98
2	1002403	THE VILLAGE OF HARLEM....NEW YORK !	78829239556	NaN	Elise	Manhattan	Harlem	40.80902	-73.94
3	1002755	NaN	85098326012	unconfirmed	Garry	Brooklyn	Clinton Hill	40.68514	-73.95
4	1003689	Entire Apt: Spacious Studio/Loft by central park	92037596077	verified	Lyndon	Manhattan	East Harlem	40.79851	-73.94

5 rows × 26 columns

Checking here how many initial rows of data there are and what type of data each column is.

```
In [4]: open_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 102599 entries, 0 to 102598
Data columns (total 26 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     102599 non-null  int64
1   NAME                                  102349 non-null  object
2   host_id                               102599 non-null  int64
3   host_identity_verified                102310 non-null  object
4   host name                             102193 non-null  object
5   neighbourhood group                  102570 non-null  object
6   neighbourhood                         102583 non-null  object
7   lat                                   102591 non-null  float64
8   long                                  102591 non-null  float64
9   country                              102067 non-null  object
10  country code                         102468 non-null  object
11  instant_bookable                     102494 non-null  object
12  cancellation_policy                  102523 non-null  object
13  room type                            102599 non-null  object
14  Construction year                    102385 non-null  float64
15  price                                102352 non-null  object
16  service fee                          102326 non-null  object
17  minimum nights                       102190 non-null  float64
18  number of reviews                    102416 non-null  float64
19  last review                          86706 non-null  object
20  reviews per month                    86720 non-null  float64
21  review rate number                   102273 non-null  float64
22  calculated host listings count        102280 non-null  float64
23  availability 365                      102151 non-null  float64
24  house_rules                          50468 non-null  object
25  license                               2 non-null      object
dtypes: float64(9), int64(2), object(15)
memory usage: 20.4+ MB
```

Could potentially use ID, neighborhood group, neighborhood, lat, long, service fee, minimum nights, number of reviews, last review, reviews per month, availability,

```
In [5]: # Warning in column 25 due to multiple dtypes so removing for now
open_data.drop(columns='license', axis=1, inplace = True)
open_data.head()
```

Out[5]:

	id	NAME	host id	host_identity_verified	host name	neighbourhood group	neighbourhood	lat	lon
0	1001254	Clean & quiet apt home by the park	80014485718	unconfirmed	Madaline	Brooklyn	Kensington	40.64749	-73.97
1	1002102	Skylit Midtown Castle	52335172823	verified	Jenna	Manhattan	Midtown	40.75362	-73.98
2	1002403	THE VILLAGE OF HARLEM....NEW YORK!	78829239556	NaN	Elise	Manhattan	Harlem	40.80902	-73.94
3	1002755	NaN	85098326012	unconfirmed	Garry	Brooklyn	Clinton Hill	40.68514	-73.95
4	1003689	Entire Apt: Spacious Studio/Loft by central park	92037596077	verified	Lyndon	Manhattan	East Harlem	40.79851	-73.94

5 rows × 25 columns

Seeing the shape of the dataset. 102599 rows by 25 columns.

```
In [6]: open_data.shape
```

Out[6]: (102599, 25)

Could find out neighborhood with most airbnbs, how long people use book for,

Checking to see if there are other countries in the dataset other than USA.

```
In [7]: # Checking to see if this is a constant value
open_data.country.unique()
```

Out[7]: array(['United States', nan], dtype=object)

Because this value is constant, it can be removed.

```
In [8]: # Removing this column since it is constant
open_data.drop(columns='country', axis=1, inplace = True)
```

Choosing the columns I think I will need going forward and calling them main\_column.

```
In [9]: # Removing host id and host name b/c it's not necessary and last review for now
main_columns = ['host_identity_verified', 'neighbourhood group', 'last review',
                'neighbourhood', 'price', 'service fee', 'minimum nights', 'number of reviews',
                'reviews per month', 'review rate number',
                'calculated host listings count', 'availability 365']
# Using these columns for Folium map later on in data analysis
#df_foliumn = ['neighbourhood', 'price', 'lat', 'long' ]
```

Doing one more check to make sure I selected the correct columns.

```
In [10]: data = open_data[main_columns]
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 102599 entries, 0 to 102598
Data columns (total 12 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   host_identity_verified                102310 non-null object
 1   neighbourhood group                  102570 non-null object
 2   last review                          86706 non-null object
 3   neighbourhood                        102583 non-null object
 4   price                               102352 non-null object
 5   service fee                          102326 non-null object
 6   minimum nights                      102190 non-null float64
 7   number of reviews                   102416 non-null float64
 8   reviews per month                   86720 non-null float64
 9   review rate number                   102273 non-null float64
10   calculated host listings count       102280 non-null float64
11   availability 365                     102151 non-null float64
dtypes: float64(6), object(6)
memory usage: 9.4+ MB
```

### 3 Cleaning Data

Here is where I go through every column that has NaNs and either drop those rows or replace them with another value.

```
In [11]: # Checking for NaNs
data.isna().sum()
```

```
Out[11]: host_identity_verified                289
neighbourhood group                  29
last review                          15893
neighbourhood                        16
price                               247
service fee                          273
minimum nights                      409
number of reviews                   183
reviews per month                   15879
review rate number                   326
calculated host listings count       319
availability 365                     448
dtype: int64
```

```
In [12]: data['host_identity_verified'].unique()
```

```
Out[12]: array(['unconfirmed', 'verified', nan], dtype=object)
```

```
In [13]: # Switching host_identity_verified's NaNs to unconfirmed
data['host_identity_verified'] = data['host_identity_verified'].fillna('unconfirmed')
data['host_identity_verified'].isna().sum()
#data.isna().sum()
```

Out[13]: 0

```
In [14]: # Changing NaNs of host name to Unknown
#data['host name'] = data['host name'].fillna('Unknown')
#data.isna().sum()
```

These rows get dropped because 29 and 16 rows is insignificant when there is a total of 100k+.

```
In [15]: # Dropping the 29 NaNs from neighborhood group
data.dropna(subset= ['neighbourhood group', 'neighbourhood'], inplace=True)
data.isna().sum()
```

```
Out[15]: host_identity_verified      0
neighbourhood group                0
last review                      15889
neighbourhood                    0
price                           245
service fee                      273
minimum nights                   407
number of reviews                183
reviews per month                15877
review rate number               324
calculated host listings count   319
availability 365                 436
dtype: int64
```

```
In [16]: data['neighbourhood group'].unique()
```

```
Out[16]: array(['Brooklyn', 'Manhattan', 'brookln', 'manhatan', 'Queens',
                'Staten Island', 'Bronx'], dtype=object)
```

Here I noticed Brooklyn and Manhattan were spelled wrong so those names needed to be replaced.

```
In [17]: # Replacing brooklyn with Brooklyn and manhatan with Manhattan
data.loc[:, ['neighbourhood group']] = data['neighbourhood group'].str.replace('brookln', 'B
data.loc[:, ['neighbourhood group']] = data['neighbourhood group'].str.replace('manhatan', 'M
```

```
In [18]: # Turning NaNs into $ for now
#data['price'] = data['price'].fillna('0')
```

Rows where the price was a NaN got replaced with a 0. Using a string here would over complicate things. Also 245 NaNs were insignificant.

```
In [19]: data.loc[:, ['price']] = data['price'].fillna('0')
```

One other thing wrong with the price values were that they were strings with commas and dollar signs in them. These needed to be removed so the value could be converted to an integer.

```
In [20]: # Replacing NaNs in Price column with the average price of an Airbnb in NY
# First need to convert str to integer
data.loc[:, ['price']] = data['price'].str.replace(',', '')
data.loc[:, ['price']] = data['price'].str.replace('$', '')
data.loc[:, ['price']] = data['price'].astype(int)
data['price'].dtype
```

```
Out[20]: dtype('int64')
```

```
In [21]: # Now they are all integers we need to turn all the 0's into the price avg
#data.loc[:, ['price']] = data['price'].replace(to_replace = 0, value = data['price'].mean())
#data.isna().sum()
```

Doing the same thing here for service fee.

```
In [22]: # Replacing NaNs of Service fee
data.loc[:, ['service fee']] = data['service fee'].fillna('0')
```

```
In [23]: # Doing the same thing with service fee
data.loc[:, ['service fee']] = data['service fee'].str.replace(',', '')
data.loc[:, ['service fee']] = data['service fee'].str.replace('$', '')
data.loc[:, ['service fee']] = data['service fee'].astype(int)
data['service fee'].dtype
```

```
Out[23]: dtype('int64')
```

```
In [24]: data.isna().sum()
```

```
Out[24]: host_identity_verified      0
neighbourhood group                 0
last review                        15889
neighbourhood                       0
price                               0
service fee                         0
minimum nights                     407
number of reviews                   183
reviews per month                   15877
review rate number                   324
calculated host listings count      319
availability 365                    436
dtype: int64
```

Replacing minimum nights NaNs with 0's as well.

```
In [25]: data.loc[:, ['minimum nights']] = data['minimum nights'].fillna(0).astype(int)
```

```
In [26]: # Convert rest of strings into integers
data.loc[:, ['minimum nights']] = data['minimum nights'].astype(int)
```

From here, I thought it would be better to use the minimum night's average amount instead of 0. So I replaced the 0's with the min nights average.

```
In [27]: # Replace temp 0's with minimum night avg
data.loc[:, ['minimum nights']] = data.loc[:, ['minimum nights']].replace(to_replace=0, value
```

```
In [28]: data.isna().sum()
```

```
Out[28]: host_identity_verified      0
          neighbourhood group        0
          last review                15889
          neighbourhood              0
          price                      0
          service fee                0
          minimum nights             0
          number of reviews          183
          reviews per month          15877
          review rate number         324
          calculated host listings count 319
          availability 365            436
          dtype: int64
```

Decided to do the same thing with the number of reviews. I was trying to keep as many rows as possible.

```
In [29]: # Replace NaNs with # of Review avg
         data.loc[:, ['number of reviews']] = data.loc[:, ['number of reviews']].replace(to_replace =
```

```
In [30]: # Convert floats to Integers
         data.loc[:, ['number of reviews']] = data.loc[:, ['number of reviews']].astype(int)
         data.isna().sum()
```

```
Out[30]: host_identity_verified      0
          neighbourhood group        0
          last review                15889
          neighbourhood              0
          price                      0
          service fee                0
          minimum nights             0
          number of reviews          0
          reviews per month          15877
          review rate number         324
          calculated host listings count 319
          availability 365            436
          dtype: int64
```

There was no simple way to convert the last review dates from NaNs into something meaningful, so unfortunately they needed to be dropped.

```
In [31]: # Replace NaNs with N/A because there is too much useful data to drop
         data.dropna(subset=['last review'], inplace=True)
         # Using N/A in place of the NaNs
         #data['last review'] = data['last review'].fillna('N/A')
         data.isna().sum()
```

```
Out[31]: host_identity_verified      0
          neighbourhood group        0
          last review                0
          neighbourhood              0
          price                      0
          service fee                0
          minimum nights             0
          number of reviews          0
          reviews per month          12
          review rate number         287
          calculated host listings count 261
          availability 365            170
          dtype: int64
```

For future analysis, the remaining values for last review needed to first be converted to a datetime type.



```
In [32]: # Converting last review column to datetime dtype

data['last review'] = pd.to_datetime(data['last review'])
```

```
In [33]: data['last review']
```

```
Out[33]: 0      2021-10-19
1      2022-05-21
3      2019-07-05
4      2018-11-19
5      2019-06-22
...
102588 2019-06-29
102591 2019-01-04
102593 2015-09-06
102595 2015-07-06
102597 2015-10-11
Name: last review, Length: 86665, dtype: datetime64[ns]
```

Because our models require our values to be either categorical or continuous, year, month and day needed to be in their own column. So that is what this code is doing.

```
In [34]: # Create 3 new columns for year, month, and day
#Year
data['lr_year'] = pd.DatetimeIndex(data['last review']).year
#Month
data['lr_month'] = pd.DatetimeIndex(data['last review']).month
#Day
data['lr_day'] = pd.DatetimeIndex(data['last review']).day
```

```
In [35]: # No need for the last review column anymore so dropping it
data.drop(columns='last review', axis=1, inplace=True)
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 86665 entries, 0 to 102597
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                ---
0   host_identity_verified                86665 non-null  object
1   neighbourhood group                  86665 non-null  object
2   neighbourhood                        86665 non-null  object
3   price                                86665 non-null  int64
4   service fee                          86665 non-null  int64
5   minimum nights                      86665 non-null  float64
6   number of reviews                   86665 non-null  int64
7   reviews per month                   86653 non-null  float64
8   review rate number                   86378 non-null  float64
9   calculated host listings count       86404 non-null  float64
10  availability 365                     86495 non-null  float64
11  lr_year                             86665 non-null  int64
12  lr_month                            86665 non-null  int64
13  lr_day                              86665 non-null  int64
dtypes: float64(5), int64(6), object(3)
memory usage: 9.9+ MB
```

```
In [36]: # Replace NaNs with average of reviews per month because there is too much useful data to dr
#data['reviews per month'] = data['reviews per month'].fillna('N/A')
data.loc[:, ['reviews per month']] = data.loc[:, ['reviews per month']].replace(to_replace =
data.isna().sum())
```

```
Out[36]: host_identity_verified      0
neighbourhood group                0
neighbourhood                      0
price                             0
service fee                        0
minimum nights                    0
number of reviews                 0
reviews per month                  0
review rate number                 287
calculated host listings count     261
availability 365                   170
lr_year                           0
lr_month                          0
lr_day                            0
dtype: int64
```

```
In [37]: # Convert last review to datetime
# data['last review'] = pd.to_datetime(data['last review'])
```

The remaining NaNs were too small to care about so they were also dropped.

```
In [38]: # Dropping NaNs of Review rate number, calculated host listings, and availability
data.dropna(subset= ['review rate number', 'calculated host listings count', 'availability 3
data.isna().sum())
```

```
Out[38]: host_identity_verified      0
neighbourhood group                0
neighbourhood                      0
price                             0
service fee                        0
minimum nights                    0
number of reviews                 0
reviews per month                  0
review rate number                 0
calculated host listings count     0
availability 365                   0
lr_year                           0
lr_month                          0
lr_day                            0
dtype: int64
```

In [39]:

data.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 86017 entries, 0 to 102597
Data columns (total 14 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   host_identity_verified                86017 non-null  object
 1   neighbourhood group                   86017 non-null  object
 2   neighbourhood                         86017 non-null  object
 3   price                                86017 non-null  int64
 4   service fee                          86017 non-null  int64
 5   minimum nights                       86017 non-null  float64
 6   number of reviews                    86017 non-null  int64
 7   reviews per month                    86017 non-null  float64
 8   review rate number                   86017 non-null  float64
 9   calculated host listings count       86017 non-null  float64
10   availability 365                     86017 non-null  float64
11   lr_year                              86017 non-null  int64
12   lr_month                             86017 non-null  int64
13   lr_day                               86017 non-null  int64
dtypes: float64(5), int64(6), object(3)
memory usage: 9.8+ MB
```

Finally I have a dataset with cleaned data. Here is where I take a look at the final product.

In [40]:

data.head()

Out[40]:

	host_identity_verified	neighbourhood group	neighbourhood	price	service fee	minimum nights	number of reviews	reviews per month	review rate number	calculated host listings count
0	unconfirmed	Brooklyn	Kensington	966	193	10.0	9	0.21	4.0	6.0
1	verified	Manhattan	Midtown	142	28	30.0	45	0.38	4.0	2.0
3	unconfirmed	Brooklyn	Clinton Hill	368	74	30.0	270	4.64	4.0	1.0
4	verified	Manhattan	East Harlem	204	41	10.0	9	0.10	3.0	1.0
5	verified	Manhattan	Murray Hill	577	115	3.0	74	0.59	3.0	1.0

In [41]:

data.describe()

Out[41]:

	price	service fee	minimum nights	number of reviews	reviews per month	review rate number	calculated host listings count	availability 365
count	86017.000000	86017.000000	86017.000000	86017.000000	86017.000000	86017.000000	86017.000000	86017.000000
mean	624.511922	124.830336	7.471016	32.253183	1.373066	3.279329	7.030122	141.920225
std	332.793549	66.575394	27.907669	51.883511	1.744690	1.283225	29.380053	133.910756
min	0.000000	0.000000	-365.000000	1.000000	0.010000	1.000000	1.000000	-10.000000
25%	338.000000	68.000000	2.000000	3.000000	0.220000	2.000000	1.000000	6.000000
50%	623.000000	125.000000	3.000000	11.000000	0.740000	3.000000	1.000000	101.000000
75%	913.000000	183.000000	5.000000	38.000000	2.000000	4.000000	2.000000	266.000000
max	1200.000000	240.000000	5645.000000	1024.000000	90.000000	5.000000	332.000000	3677.000000

▼

3.1 Checking Value Counts of Columns

There are over 10,000 columns when data is one hot encoded, so this number needs to be reduced.

```
In [42]: data['host_identity_verified'].value_counts()
```

```
Out[42]: unconfirmed    43147
verified      42870
Name: host_identity_verified, dtype: int64
```

This just has 2 different values.

```
In [43]: #data['host name'].value_counts(normalize=True).to_frame().style.format('{:.2%}')
```

Has 13131 different hosts. To reduce this amount, hosts with only 5% of the total value counts will be categorized as "Other".

I later realized the Host name values did not matter in the overall analysis of the models. So they were removed entirely.

```
In [44]: #data.loc[:, ['reviews per month']]
#data['host name'].value_counts(normalize=True).mul(100).round(1).astype(str) + '%'
```

Checking on which was the most popular boroughs.

```
In [45]: data['neighbourhood group'].value_counts()
```

```
Out[45]: Brooklyn      35699
Manhattan      35685
Queens         11445
Bronx          2340
Staten Island   848
Name: neighbourhood group, dtype: int64
```

There are only 5 neighborhood groups.

```
In [46]: data['neighbourhood'].value_counts()
```

```
Out[46]: Bedford-Stuyvesant    6963
Williamsburg      6526
Harlem            4712
Bushwick          4102
Hell's Kitchen     3275
...
Gerritsen Beach      3
Glen Oaks            2
Woodrow             1
New Dorp            1
Chelsea, Staten Island 1
Name: neighbourhood, Length: 223, dtype: int64
```

This is where things became a little difficult. There were too many neighborhoods, which caused our models later on to take too long. 223 neighborhoods to be exact causing 223 extra columns to be created.

```
In [47]: data['neighbourhood'].value_counts(normalize=True)
```

```
Out[47]: Bedford-Stuyvesant      0.080949
Williamsburg      0.075869
Harlem            0.054780
Bushwick         0.047688
Hell's Kitchen    0.038074
...
Gerritsen Beach   0.000035
Glen Oaks         0.000023
Woodrow           0.000012
New Dorp          0.000012
Chelsea, Staten Island 0.000012
Name: neighbourhood, Length: 223, dtype: float64
```

```
In [48]: data['neighbourhood'].value_counts(normalize=True).mul(100).round(1).astype(str) + '%'
```

```
Out[48]: Bedford-Stuyvesant      8.1%
Williamsburg      7.6%
Harlem            5.5%
Bushwick         4.8%
Hell's Kitchen    3.8%
...
Gerritsen Beach   0.0%
Glen Oaks         0.0%
Woodrow           0.0%
New Dorp          0.0%
Chelsea, Staten Island 0.0%
Name: neighbourhood, Length: 223, dtype: object
```

The value counts number of each neighborhood needed to be first converted to a list.

```
In [49]: # Storing value counts in a list
#neighborhood_vc = list(data['neighbourhood'].value_counts(normalize=True).mul(100).round(2))
#data['neighbourhood'].value_counts(normalize=True).mul(100).round(2).values
neighborhood_vc = list(data['neighbourhood'].value_counts())
neighborhood_vc
```

```
Out[49]: [6963,
6526,
4712,
4102,
3275,
3101,
3031,
2891,
2774,
2496,
2047,
1832,
1721,
1654,
1565,
1470,
1229,
1137,
1087,
1000]
```

There are 223 different neighborhoods.

To reduce the amount of neighborhoods, I need to only include neighborhoods with a value counts of over 1000.

```
In [50]: # Get list of value counts greater than 1000
new_neighborhood = [x for x in neighborhood_vc if x > 1000]
#for neighborhood in neighborhood_vc:
#    if neighborhood_vc < 10:
#        other = other + neighborhood_vc[neighborhood]

#print('Sum percentage neighborhoods that are less than 1 percent of the data ', other)
print(new_neighborhood)
```

```
[6963, 6526, 4712, 4102, 3275, 3101, 3031, 2891, 2774, 2496, 2047, 1832, 1721, 1654, 1565,
1470, 1229, 1137, 1087, 1083, 1022]
```

```
In [51]: # This is the sum of all other neighborhoods that people stayed less than 1000 times
other = sum(neighborhood_vc) - sum(new_neighborhood)
other
```

Out[51]: 30299

After removing neighborhoods visited less than 1000 times, I was left with 21 neighborhoods.

```
In [52]: len(new_neighborhood)
```

Out[52]: 21

```
In [53]: # The Top 21 most visited neighborhoods with more than 1000 visits
top_21 = data['neighbourhood'].value_counts().index[:21].tolist()
top_21
```

```
Out[53]: ['Bedford-Stuyvesant',
'Williamsburg',
'Harlem',
'Bushwick',
"Hell's Kitchen",
'Upper West Side',
'Upper East Side',
'East Village',
'Crown Heights',
'Midtown',
'East Harlem',
'Chelsea',
'Greenpoint',
'Lower East Side',
'Astoria',
'Washington Heights',
'West Village',
'Financial District',
'Flatbush',
'East Flatbush',
'Prospect-Lefferts Gardens']
```

Neighborhoods outside of the top 21 needed to be converted to 'Other'.

```
In [54]: # Get list of neighborhoods outside top 21
outside_21 = data['neighbourhood'].value_counts().index[21:].tolist()
outside_21
```

```
Out[54]: ['Clinton Hill',
'Long Island City',
'Flushing',
'Park Slope',
'Fort Greene',
'Sunset Park',
'Kips Bay',
'Ridgewood',
'Chinatown',
'Murray Hill',
'Sunnyside',
'Greenwich Village',
'Prospect Heights',
'Ditmars Steinway',
'SoHo',
'Gramercy',
'East New York',
'Jamaica',
'South Slope',
'Woodside',
'Morningside Heights',
'Inwood',
'Elmhurst',
'East Elmhurst',
'Nolita',
'Gowanus',
'Carroll Gardens',
'Jackson Heights',
'Theater District',
'Canarsie',
'Cypress Hills',
'Kensington',
'Boerum Hill',
'Sheepshead Bay',
'Windsor Terrace',
'Bay Ridge',
'Forest Hills',
'Brooklyn Heights',
'Tribeca',
'Springfield Gardens',
'Little Italy',
'Borough Park',
'Maspeth',
'Rego Park',
'Arverne',
'St. Albans',
'Richmond Hill',
'Midwood',
'Flatlands',
'Cobble Hill',
'Rockaway Beach',
'Woodhaven',
'Red Hook',
'Rosedale',
'Corona',
'Brownsville',
'Brighton Beach',
'Wakefield',
'Roosevelt Island',
'Bensonhurst',
'Mott Haven',
'Kingsbridge',
'Fordham',
'Downtown Brooklyn',
```

'Flatiron District',  
'Ozone Park',  
'South Ozone Park',  
'Glendale',  
'Queens Village',  
'NoHo',  
'Two Bridges',  
'Bayside',  
'Longwood',  
'Concourse',  
'Gravesend',  
'St. George',  
'Fort Hamilton',  
'Tompkinsville',  
'Briarwood',  
'Fresh Meadows',  
'Parkchester',  
'Port Morris',  
'Williamsbridge',  
'Allerton',  
'Columbia St',  
'Pelham Gardens',  
'Kew Gardens',  
'Civic Center',  
'Battery Park City',  
'Far Rockaway',  
'Laurelton',  
'Highbridge',  
'Concourse Village',  
'Norwood',  
'Stuyvesant Town',  
'Middle Village',  
'Cambria Heights',  
'Stapleton',  
'Vinegar Hill',  
'Concord',  
'Kew Gardens Hills',  
'West Brighton',  
'Throgs Neck',  
'Mount Hope',  
'DUMBO',  
'Arrochar',  
'Hunts Point',  
'Coney Island',  
'Schuylerville',  
'Bath Beach',  
'University Heights',  
'Clason Point',  
'Bronxdale',  
'Claremont Village',  
'Howard Beach',  
'Jamaica Estates',  
'College Point',  
'Morris Park',  
'Belmont',  
'Pelham Bay',  
'Randall Manor',  
'City Island',  
'Clifton',  
'Morris Heights',  
'Hollis',  
'Soundview',  
'Bayswater',  
'Tremont',  
'Woodlawn',  
'Baychester',  
'Bergen Beach',  
'Edenwald',



```
'Navy Yard',
'Belle Harbor',
'Edgemere',
'Fieldston',
'Great Kills',
'Marble Hill',
'Dyker Heights',
'Eastchester',
'Van Nest',
'Morrisania',
'Mariners Harbor',
'Douglaston',
'Westchester Square',
'Port Richmond',
'Manhattan Beach',
'Grant City',
'East Morrisania',
'Jamaica Hills',
'Melrose',
'North Riverdale',
'Unionport',
'Whitestone',
'Bellerose',
'Grymes Hill',
'Olinville',
'Mount Eden',
'Shore Acres',
'South Beach',
'New Springville',
'New Brighton',
'Midland Beach',
'Rosebank',
'Oakwood',
'Sea Gate',
'Mill Basin',
'Riverdale',
'Dongan Hills',
'Castleton Corners',
'Tottenville',
'Arden Heights',
'Huguenot',
'Silver Lake',
'Spuyten Duyvil',
'Castle Hill',
'Todt Hill',
'New Dorp Beach',
'Bay Terrace',
'Breezy Point',
'Neponsit',
'Graniteville',
'Emerson Hill',
'West Farms',
'Little Neck',
'Howland Hook',
'Eltingville',
'Holliswood',
'Bull's Head',
'Westerleigh',
'Richmondtown',
'Co-op City',
'Prince's Bay',
'Bay Terrace, Staten Island',
'Lighthouse Hill',
'Rossville',
'Willowbrook',
'Gerritsen Beach',
'Glen Oaks',
'Woodrow',
```

```
'New Dorp',
'Chelsea, Staten Island']
```

Creating a DF with just neighborhoods in the top 21.

```
In [55]: # Dropping rows that contain neighbors outside of top 21
top21_df = data[data['neighbourhood'].isin(top_21)]
top21_df
```

Out[55]:

	host_identity_verified	neighbourhood group	neighbourhood	price	service fee	minimum nights	number of reviews	reviews per month	review rate number	calculate ho listing cou
1	verified	Manhattan	Midtown	142	28	30.0	45	0.38	4.0	2
4	verified	Manhattan	East Harlem	204	41	10.0	9	0.10	3.0	1
6	unconfirmed	Brooklyn	Bedford-Stuyvesant	71	14	45.0	49	0.40	5.0	1
7	unconfirmed	Brooklyn	Bedford-Stuyvesant	1060	212	45.0	49	0.40	5.0	1
8	verified	Manhattan	Hell's Kitchen	1018	204	2.0	430	3.47	3.0	1
...	...	...	...	...	...	...	...	...	...	...
102585	verified	Brooklyn	Williamsburg	643	129	5.0	4	0.08	5.0	1
102586	verified	Manhattan	Upper East Side	208	42	4.0	8	0.17	5.0	1

Here is where the neighborhoods outside of top 21 got converted to Other.

```
In [56]: # If neighborhood is outside of top 21 change neighborhood name to 'Other'
outside21_df = data[data['neighbourhood'].isin(outside_21)]
outside21_df
```

Out[56]:

	host_identity_verified	neighbourhood group	neighbourhood	price	service fee	minimum nights	number of reviews	reviews per month	review rate number	calcula h listir co
0	unconfirmed	Brooklyn	Kensington	966	193	10.0	9	0.21	4.0	
3	unconfirmed	Brooklyn	Clinton Hill	368	74	30.0	270	4.64	4.0	
5	verified	Manhattan	Murray Hill	577	115	3.0	74	0.59	3.0	
10	verified	Manhattan	Chinatown	319	64	1.0	160	1.33	3.0	
13	verified	Brooklyn	South Slope	580	116	4.0	167	1.34	4.0	
...	...	...	...	...	...	...	...	...	...	...
102580	verified	Queens	East Elmhurst	609	122	1.0	209	4.38	3.0	
102584	verified	Queens	Arverne	566	113	2.0	89	1.82	2.0	
102588	unconfirmed	Manhattan	Flatiron District	618	124	1.0	177	3.78	4.0	
102595	unconfirmed	Manhattan	Morningside Heights	837	167	1.0	1	0.02	2.0	
102597	unconfirmed	Queens	Long Island City	546	109	2.0	5	0.10	3.0	

30299 rows × 14 columns

```
In [57]: # Set the neighborhood groups outside of top 21 to Other
outside21_df.loc[:, 'neighbourhood'] = 'Other'
outside21_df
```

Out[57]:

	host_identity_verified	neighbourhood group	neighbourhood	price	service fee	minimum nights	number of reviews	reviews per month	review rate number	calculated listing cost
0	unconfirmed	Brooklyn	Other	966	193	10.0	9	0.21	4.0	
3	unconfirmed	Brooklyn	Other	368	74	30.0	270	4.64	4.0	
5	verified	Manhattan	Other	577	115	3.0	74	0.59	3.0	
10	verified	Manhattan	Other	319	64	1.0	160	1.33	3.0	
13	verified	Brooklyn	Other	580	116	4.0	167	1.34	4.0	
...	...	...	...	...	...	...	...	...	...	
102580	verified	Queens	Other	609	122	1.0	209	4.38	3.0	
102584	verified	Queens	Other	566	113	2.0	89	1.82	2.0	
102588	unconfirmed	Manhattan	Other	618	124	1.0	177	3.78	4.0	
102595	unconfirmed	Manhattan	Other	837	167	1.0	1	0.02	2.0	
102597	unconfirmed	Queens	Other	546	109	2.0	5	0.10	3.0	

30299 rows × 14 columns

Now the top 21 DF and the newly converted outside 21 DF get combined back together.

```
In [58]: # Concat back top 21 df and outside 21 df
combined_df = [top21_df, outside21_df]
main_df = pd.concat(combined_df)
main_df
```

Out[58]:

	host_identity_verified	neighbourhood group	neighbourhood	price	service fee	minimum nights	number of reviews	reviews per month	review rate number	calculated listing cost
1	verified	Manhattan	Midtown	142	28	30.0	45	0.38	4.0	
4	verified	Manhattan	East Harlem	204	41	10.0	9	0.10	3.0	
6	unconfirmed	Brooklyn	Bedford-Stuyvesant	71	14	45.0	49	0.40	5.0	
7	unconfirmed	Brooklyn	Bedford-Stuyvesant	1060	212	45.0	49	0.40	5.0	
8	verified	Manhattan	Hell's Kitchen	1018	204	2.0	430	3.47	3.0	
...	...	...	...	...	...	...	...	...	...	
102580	verified	Queens	Other	609	122	1.0	209	4.38	3.0	
102584	verified	Queens	Other	566	113	2.0	89	1.82	2.0	
102588	unconfirmed	Manhattan	Other	618	124	1.0	177	3.78	4.0	
102595	unconfirmed	Manhattan	Other	837	167	1.0	1	0.02	2.0	
102597	unconfirmed	Queens	Other	546	109	2.0	5	0.10	3.0	

86017 rows × 14 columns

One more check of the value counts on neighborhood and now I have a more manageable amount of neighborhoods.

```
In [59]: # Check back value counts
#main_df[main_df['neighbourhood'] == 'Other']
main_df['neighbourhood'].value_counts()
```

```
Out[59]: Other                30299
Bedford-Stuyvesant          6963
Williamsburg                6526
Harlem                      4712
Bushwick                    4102
Hell's Kitchen              3275
Upper West Side             3101
Upper East Side             3031
East Village                2891
Crown Heights               2774
Midtown                     2496
East Harlem                 2047
Chelsea                     1832
Greenpoint                  1721
Lower East Side             1654
Astoria                     1565
Washington Heights          1470
West Village                1229
Financial District          1137
Flatbush                    1087
East Flatbush               1083
Prospect-Lefferts Gardens   1022
Name: neighbourhood, dtype: int64
```

Checking range of years people booked for.

```
In [60]: main_df['lr_year'].value_counts()
```

```
Out[60]: 2019      42558
2018      11378
2022      10282
2021       6705
2017       6543
2016       4282
2020       2056
2015       1860
2014        244
2013         80
2012         26
2040          1
2026          1
2058          1
Name: lr_year, dtype: int64
```

Bookings range from 2012 to 2022.




## 4 Data Analysis

```
In [61]: neighborhood_count = main_df['neighbourhood'].value_counts().head(25)
neighborhood_count
```

```
Out[61]: Other                30299
Bedford-Stuyvesant          6963
Williamsburg                6526
Harlem                      4712
Bushwick                   4102
Hell's Kitchen              3275
Upper West Side            3101
Upper East Side            3031
East Village               2891
Crown Heights              2774
Midtown                    2496
East Harlem                2047
Chelsea                    1832
Greenpoint                 1721
Lower East Side            1654
Astoria                    1565
Washington Heights         1470
West Village               1229
Financial District          1137
Flatbush                   1087
East Flatbush              1083
Prospect-Lefferts Gardens  1022
Name: neighbourhood, dtype: int64
```

```
In [62]: neighborhood_count.plot(kind='bar',figsize=(15,6))
plt.ylabel('Number of AirBnBs')
plt.xlabel('Neighborhoods')
plt.title('AirBnBs by Neighbourhood');
```



Seems like people frequently stay in Bedford-Stuyvesant and Williamsburg the most.


```
In [63]: # Neighborhood Average Price
neighborhood_price = main_df[['price', 'neighbourhood', 'minimum nights']].groupby('neighbourh
neighborhood_price
```

Out[63]:

	price	minimum nights
neighbourhood		
Flatbush	649.732291	5.634775
West Village	639.384052	8.522622
East Flatbush	638.896584	7.335646
Prospect-Lefferts Gardens	637.556751	5.952252
Astoria	636.775719	6.466647
Bushwick	629.920770	7.470310
Other	627.626126	7.092746
Harlem	626.583829	7.300790
Lower East Side	624.380895	6.172859
Williamsburg	623.298192	6.774705
Chelsea	623.123362	9.607268
Bedford-Stuyvesant	622.865862	6.296154
Crown Heights	622.171233	6.798886
East Harlem	621.854421	6.445140
Upper West Side	618.917768	9.845602
Financial District	618.442392	7.967902
East Village	618.354549	7.200004
Hell's Kitchen	618.325802	9.078994
Upper East Side	617.705048	11.580703
Greenpoint	615.894829	7.782978
Midtown	604.850561	9.650204
Washington Heights	597.857823	7.688236

Getting a look at which neighborhoods are most expensive on average.

```
In [64]: neighborhood_price.plot(kind='bar',figsize=(18,10))  
plt.ylabel('Average Price ' )  
plt.title('Average Price of AirBnBs by Neighbourhood');
```



People spend the most money in Flatbush, West Village, East Flatbush, Prospect-Lefferts Gardens, and Astoria.

```
In [77]: # Neighborhood Average Price sorted by minimum nights
neighborhood_min_nights = main_df[['neighbourhood', 'minimum nights']].groupby('neighbourhood')
neighborhood_min_nights
```


Out[77]:

minimum nights	
neighbourhood	
Upper East Side	11.580703
Upper West Side	9.845602
Midtown	9.650204
Chelsea	9.607268
Hell's Kitchen	9.078994
West Village	8.522622
Financial District	7.967902
Greenpoint	7.782978
Washington Heights	7.688236
Bushwick	7.470310

Sorting neighborhoods by most nights spent.



```
In [79]: neighborhood_min_nights.plot(kind='bar',figsize=(10,9))
plt.ylabel('Minimum Nights ')
plt.title('Minimum Nights at AirBnB by Neighbourhood');
```



People stay the longest in Upper East Side, Upper West Side, Midtown, Chelsea, and Hell's Kitchen.

```
In [67]: neighborhood_group_price = data[['price','neighbourhood group', 'minimum nights']].groupby('neighbourhood_group_price')
```

Out[67]:

	price	minimum nights
<b>neighbourhood group</b>		
<b>Queens</b>	629.047969	6.070873
<b>Bronx</b>	628.259829	5.362738
<b>Brooklyn</b>	626.066669	6.836214
<b>Manhattan</b>	621.333417	8.750596
<b>Staten Island</b>	621.253538	5.062976

All of NY's boroughs seem to have about the same average price.

```
In [68]: # Creating a DF with just Neighborhood, Latitude, longitude and Price
#neighborhood_locations = open_data[df_foliumn]
#neighborhood_locations.info()
```


```
In [69]: #Converting price and removing NaNs
#neighborhood_locations.loc[:, ['price']] = neighborhood_locations['price'].fillna('0')

# Replacing NaNs in Price column with the average price of an Airbnb in NY
# First need to convert str to integer
#neighborhood_locations.loc[:, ['price']] = neighborhood_locations['price'].str.replace(',', '')
#neighborhood_locations.loc[:, ['price']] = neighborhood_locations['price'].str.replace('$', '')
#neighborhood_locations.loc[:, ['price']] = neighborhood_locations['price'].astype(int)
#neighborhood_locations['price'].dtype
```

```
In [70]: #Check for NaNs
#neighborhood_locations.isna().sum()
```

```
In [71]: #neighborhood_locations.dropna(subset= ['neighbourhood', 'lat', 'long'], inplace=True)
# Final Check for NaNs
#neighborhood_locations.isna().sum()
```


```
In [72]: data.hist(figsize = (20,18));
```



Checking to see if any variable was normally distributed.

Not much normal distribution going on within these graphs.

```
In [73]: # Relationship between price and service fee
plt.figure(figsize=(15,10))
plt.title("Relationship between Price and Service Fee", size=25, color="red")
sns.scatterplot(x=data.price, y=data['service fee'] );
```

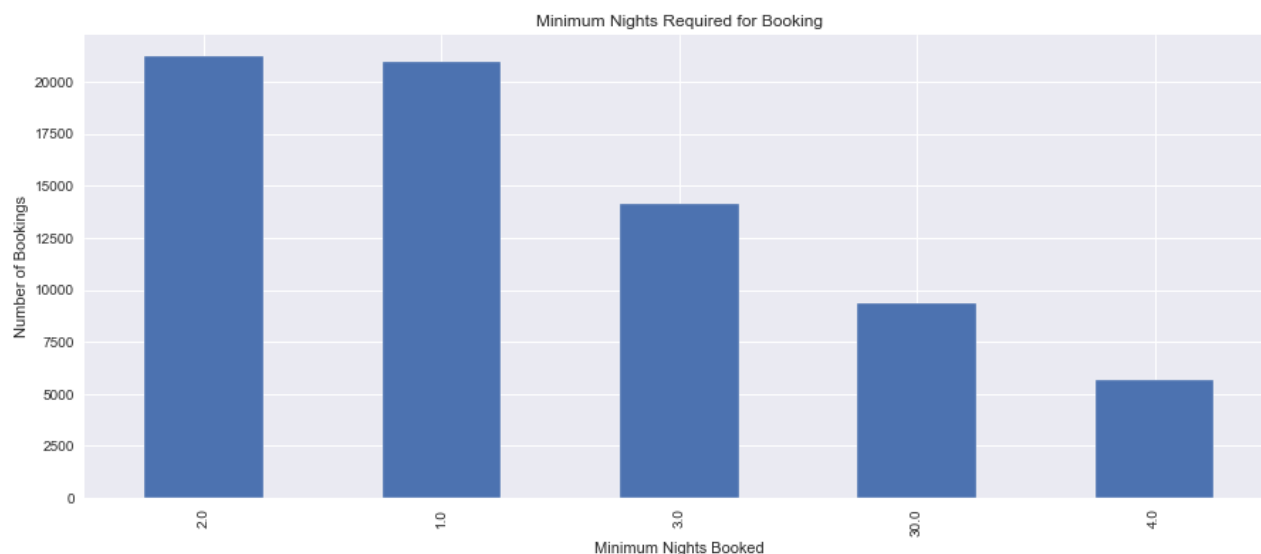


High room price means the service price will also be high.

```
In [75]: min_nights = data['minimum nights'].value_counts().head()
min_nights
```

```
Out[75]: 2.0      21193
         1.0      20963
         3.0      14143
        30.0       9366
         4.0       5667
         Name: minimum nights, dtype: int64
```

```
In [76]: min_nights.plot(kind='bar', figsize=(15,6))
plt.ylabel('Number of Bookings')
plt.xlabel('Minimum Nights Booked')
plt.title('Minimum Nights Required for Booking');
```



Visualizing minimum nights required for booking. Most popular is 2 nights. Closely followed is 1 night then 3. Weirdly, 4th is 30 nights and last is 4 nights. The New York City Airbnb law says it's illegal to rent apartments for fewer than 30 days in most buildings, particularly Class A dwellings.

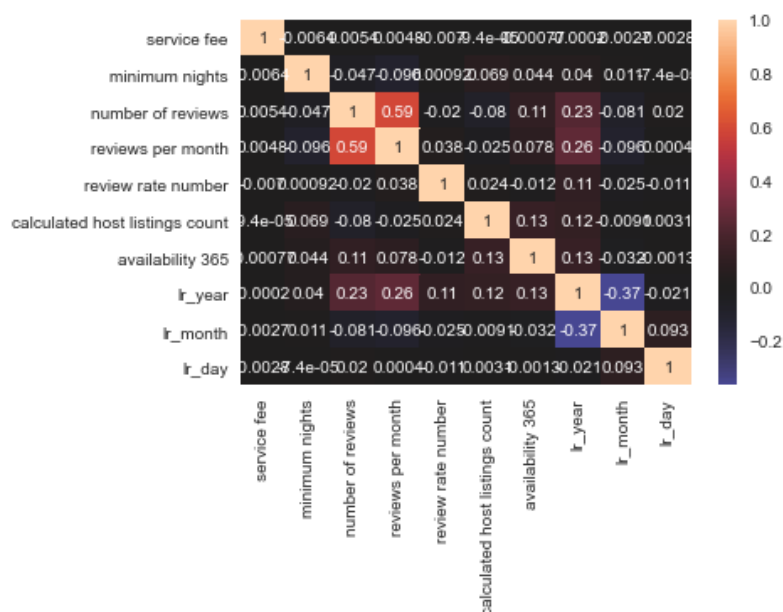
```
In [314]: data.columns
```

```
Out[314]: Index(['host_identity_verified', 'neighbourhood group', 'neighbourhood',
               'price', 'service fee', 'minimum nights', 'number of reviews',
               'reviews per month', 'review rate number',
               'calculated host listings count', 'availability 365', 'lr_year',
               'lr_month', 'lr_day'],
              dtype='object')
```

## ▼ 4.1 Check Correlation

```
In [315]: check_corr = data.drop('price', axis=1)
corr = check_corr.corr()
```

```
In [316]: sns.heatmap(corr, center=0, annot=True);
```



No columns really show any correlations to one another, so there is no need to remove and similar columns. There is some with reviews per month.

## 4.2 Separating Continuous and Categorical Data & Creating Target Variable

Prepping for the train test split and one hot encoding.

```
In [317]: # Removed 'price' from continuous and last review because it was a dtype datetime
# Removed host id and host name b/c it isn't necessary
continuous = ['service fee', 'minimum nights', 'number of reviews', 'reviews per month',
              'review rate number', 'calculated host listings count', 'availability 365',
              'lr_year', 'lr_month', 'lr_day']
categoricals = ['host_identity_verified', 'neighbourhood group', 'neighbourhood']
# Creating categorical without neighbourhood
categoricals2 = ['host_identity_verified', 'neighbourhood group']
# Creating Continuous without service fee b/c it is so highly correlated with price
continuous2 = ['minimum nights', 'number of reviews', 'reviews per month',
              'review rate number', 'calculated host listings count', 'availability 365',
              'lr_year', 'lr_month', 'lr_day']
# Creating Continuous with just top features
continuous3 = ['number of reviews', 'reviews per month',
              'calculated host listings count', 'availability 365',
              'lr_year', 'lr_month', 'lr_day']
```

Second continuous variable was created here to see what some models looked like with out service fee. Third continuous was created with just the top features in mind.

```
In [318]: data_cont = main_df[continuous]
data_cont2= main_df[continuous2]
data_cont3= main_df[continuous3]
data_cat = main_df[categoricals]
data_cat2 = main_df[categoricals2]
```

```
In [319]: # Create Target variable
target = main_df['price']
main_dataframe = main_df.drop('price', axis=1)
# Creating main_dataframe2 because service fee is so highly correlated
main_dataframe2 = main_df.drop(['price', 'service fee'], axis=1)
main_dataframe3 = main_df.drop(['price', 'service fee', 'neighbourhood', 'minimum nights',
                                'review rate number'], axis=1)
```

### ▼ 4.3 Train & Test Split + One Hot Encoding

```
In [320]: data_train, data_test, target_train, target_test = train_test_split(main_dataframe, target,
                                                                              test_size = 0.25, random
```

Creating second train test split with service fee removed from dataset.

```
In [321]: data_train2, data_test2, target_train2, target_test2 = train_test_split(main_dataframe2, tar
                                                                              test_size = 0.25, random
```

Creating a third train test split with just most important features.

```
In [322]: data_train3, data_test3, target_train3, target_test3 = train_test_split(main_dataframe3, tar
                                                                              test_size = 0.25, random
```

One hot encode categorical data.

```
In [323]: ohe = OneHotEncoder(categories="auto", sparse=False, handle_unknown="ignore")
ohe_train = ohe.fit_transform(data_train[categoricals])
data_train[categoricals].shape
```

```
Out[323]: (64512, 3)
```

One Hot encode without neighbourhoods feature.

```
In [324]: ohe2 = OneHotEncoder(categories="auto", sparse=False, handle_unknown="ignore")
ohe_train2 = ohe2.fit_transform(data_train3[categoricals2])
data_train3[categoricals2].shape
```

```
Out[324]: (64512, 2)
```

Turning categorical data into numerical.

```
In [325]: ohe_train.shape
```

```
Out[325]: (64512, 29)
```

```
In [326]: # Create transformed dataframe
ohe_train = pd.DataFrame(
    ohe_train,
    columns=ohe.get_feature_names(), index=data_train.index
)
# Replace categorical data with encoded data
data_train.drop(categoricals, axis=1, inplace=True)
data_train = pd.concat([ohe_train, data_train], axis=1)
```

```
In [327]: data_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 64512 entries, 61735 to 28036
Data columns (total 39 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   x0_unconfirmed                        64512 non-null  float64
1   x0_verified                          64512 non-null  float64
2   x1_Bronx                             64512 non-null  float64
3   x1_Brooklyn                         64512 non-null  float64
4   x1_Manhattan                        64512 non-null  float64
5   x1_Queens                           64512 non-null  float64
6   x1_Staten Island                    64512 non-null  float64
7   x2_Astoria                          64512 non-null  float64
8   x2_Bedford-Stuyvesant               64512 non-null  float64
9   x2_Bushwick                         64512 non-null  float64
10  x2_Chelsea                          64512 non-null  float64
11  x2_Crown Heights                   64512 non-null  float64
12  x2_East Flatbush                   64512 non-null  float64
13  x2_East Harlem                     64512 non-null  float64
14  x2_East Village                     64512 non-null  float64
15  x2_Financial District               64512 non-null  float64
16  x2_Flatbush                        64512 non-null  float64
17  x2_Greenpoint                      64512 non-null  float64
18  x2_Harlem                          64512 non-null  float64
19  x2_Hell's Kitchen                  64512 non-null  float64
20  x2_Lower East Side                 64512 non-null  float64
21  x2_Midtown                         64512 non-null  float64
22  x2_Other                           64512 non-null  float64
23  x2_Prospect-Lefferts Gardens       64512 non-null  float64
24  x2_Upper East Side                 64512 non-null  float64
25  x2_Upper West Side                 64512 non-null  float64
26  x2_Washington Heights              64512 non-null  float64
27  x2_West Village                    64512 non-null  float64
28  x2_Williamsburg                    64512 non-null  float64
29  service fee                        64512 non-null  int64
30  minimum nights                     64512 non-null  float64
31  number of reviews                  64512 non-null  int64
32  reviews per month                  64512 non-null  float64
33  review rate number                  64512 non-null  float64
34  calculated host listings count     64512 non-null  float64
35  availability 365                    64512 non-null  float64
36  lr_year                            64512 non-null  int64
37  lr_month                           64512 non-null  int64
38  lr_day                             64512 non-null  int64
dtypes: float64(34), int64(5)
memory usage: 19.7 MB
```

```
In [328]: # Merging second dataset without service fee
# Replace categorical data with encoded data
data_train2.drop(categoricals, axis=1, inplace=True)
data_train2 = pd.concat([ohe_train, data_train2], axis=1)
data_train2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 64512 entries, 61735 to 28036
Data columns (total 38 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   x0_unconfirmed                        64512 non-null  float64
 1   x0_verified                          64512 non-null  float64
 2   x1_Bronx                             64512 non-null  float64
 3   x1_Brooklyn                         64512 non-null  float64
 4   x1_Manhattan                        64512 non-null  float64
 5   x1_Queens                           64512 non-null  float64
 6   x1_Staten Island                    64512 non-null  float64
 7   x2_Astoria                          64512 non-null  float64
 8   x2_Bedford-Stuyvesant               64512 non-null  float64
 9   x2_Bushwick                         64512 non-null  float64
10  x2_Chelsea                          64512 non-null  float64
11  x2_Crown Heights                   64512 non-null  float64
12  x2_East Flatbush                   64512 non-null  float64
13  x2_East Harlem                     64512 non-null  float64
14  x2_East Village                     64512 non-null  float64
15  x2_Financial District               64512 non-null  float64
16  x2_Flatbush                        64512 non-null  float64
17  x2_Greenpoint                      64512 non-null  float64
18  x2_Harlem                          64512 non-null  float64
19  x2_Hell's Kitchen                  64512 non-null  float64
20  x2_Lower East Side                 64512 non-null  float64
21  x2_Midtown                         64512 non-null  float64
22  x2_Other                           64512 non-null  float64
23  x2_Prospect-Lefferts Gardens       64512 non-null  float64
24  x2_Upper East Side                 64512 non-null  float64
25  x2_Upper West Side                 64512 non-null  float64
26  x2_Washington Heights              64512 non-null  float64
27  x2_West Village                    64512 non-null  float64
28  x2_Williamsburg                    64512 non-null  float64
29  minimum nights                     64512 non-null  float64
30  number of reviews                  64512 non-null  int64
31  reviews per month                  64512 non-null  float64
32  review rate number                  64512 non-null  float64
33  calculated host listings count      64512 non-null  float64
34  availability 365                    64512 non-null  float64
35  lr_year                             64512 non-null  int64
36  lr_month                            64512 non-null  int64
37  lr_day                             64512 non-null  int64
dtypes: float64(34), int64(4)
memory usage: 19.2 MB
```

```
In [329]: # Create transformed dataframe
ohe_train2 = pd.DataFrame(
    ohe_train2,
    columns=ohe2.get_feature_names(), index=data_train3.index
)
# Replace categorical data with encoded data
data_train3.drop(categoricals2, axis=1, inplace=True)
# Merging just top features
data_train3 = pd.concat([ohe_train2, data_train3], axis=1)
```



```
In [331]: # Checking final dataset
data_train3.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 64512 entries, 61735 to 28036
Data columns (total 14 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   x0_unconfirmed                        64512 non-null  float64
 1   x0_verified                          64512 non-null  float64
 2   x1_Bronx                             64512 non-null  float64
 3   x1_Brooklyn                         64512 non-null  float64
 4   x1_Manhattan                        64512 non-null  float64
 5   x1_Queens                           64512 non-null  float64
 6   x1_Staten Island                    64512 non-null  float64
 7   number of reviews                   64512 non-null  int64
 8   reviews per month                   64512 non-null  float64
 9   calculated host listings count      64512 non-null  float64
10   availability 365                    64512 non-null  float64
11   lr_year                             64512 non-null  int64
12   lr_month                            64512 non-null  int64
13   lr_day                              64512 non-null  int64
dtypes: float64(10), int64(4)
memory usage: 7.4 MB
```

```
In [87]: # Use test
```

```
# Create transformed dataframe
ohe_test = ohe.transform(data_test[categoricals])

ohe_test = pd.DataFrame(
    ohe_test,
    columns=ohe.get_feature_names(), index=data_test.index
)
# Replace categorical data with encoded data
data_test.drop(categoricals, axis=1, inplace=True)
data_test = pd.concat([ohe_test, data_test], axis=1)
```

```
In [332]: # Use test
```

```
# Create transformed dataframe
ohe_test2 = ohe2.transform(data_test3[categoricals2])

ohe_test2 = pd.DataFrame(
    ohe_test2,
    columns=ohe2.get_feature_names(), index=data_test3.index
)
# Replace categorical data with encoded data
data_test3.drop(categoricals2, axis=1, inplace=True)
data_test3 = pd.concat([ohe_test2, data_test3], axis=1)
```

```
In [125]: # Doing same thing with test data without service fee
```

```
# Replace categorical data with encoded data
data_test2.drop(categoricals, axis=1, inplace=True)
data_test2 = pd.concat([ohe_test, data_test2], axis=1)
```

```
In [333]: data_test3.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 21505 entries, 91561 to 88669
Data columns (total 14 columns):
 #   Column                                  Non-Null Count  Dtype  
---  -
 0   x0_unconfirmed                         21505 non-null  float64
 1   x0_verified                           21505 non-null  float64
 2   x1_Bronx                              21505 non-null  float64
 3   x1_Brooklyn                           21505 non-null  float64
 4   x1_Manhattan                           21505 non-null  float64
 5   x1_Queens                             21505 non-null  float64
 6   x1_Staten Island                       21505 non-null  float64
 7   number of reviews                     21505 non-null  int64   
 8   reviews per month                     21505 non-null  float64
 9   calculated host listings count        21505 non-null  float64
10   availability 365                       21505 non-null  float64
11   lr_year                               21505 non-null  int64   
12   lr_month                              21505 non-null  int64   
13   lr_day                                21505 non-null  int64   
dtypes: float64(10), int64(4)
memory usage: 2.5 MB
```

## ▼ 4.4 OLS Regression Results

Seeing which variables had any great impact on Airbnb price.

```
In [88]: X = data_train
        y = target_train
```

```
In [89]: X_int = sm.add_constant(X)
model = sm.OLS(y,X_int).fit()
model.summary()
```

Out[89]: OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.978
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.978
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	8.093e+04
<b>Date:</b>	Sun, 06 Nov 2022	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	21:02:37	<b>Log-Likelihood:</b>	-3.4255e+05
<b>No. Observations:</b>	64512	<b>AIC:</b>	6.852e+05
<b>Df Residuals:</b>	64475	<b>BIC:</b>	6.855e+05
<b>Df Model:</b>	36		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	-88.5131	153.530	-0.577	0.564	-389.432	212.405
<b>x0_unconfirmed</b>	-44.2086	76.765	-0.576	0.565	-194.668	106.250
<b>x0_verified</b>	-44.3045	76.765	-0.577	0.564	-194.765	106.156
<b>x1_Bronx</b>	-16.8736	30.756	-0.549	0.583	-77.154	43.407
<b>x1_Brooklyn</b>	-17.1883	30.689	-0.560	0.575	-77.340	42.963
<b>x1_Manhattan</b>	-17.6437	30.659	-0.575	0.565	-77.736	42.449
<b>x1_Queens</b>	-17.0828	30.725	-0.556	0.578	-77.304	43.139
<b>x1_Staten Island</b>	-19.7248	30.782	-0.641	0.522	-80.058	40.609
<b>x2_Astoria</b>	-6.2975	7.112	-0.885	0.376	-20.237	7.642
<b>x2_Bedford-Stuyvesant</b>	-4.0470	7.032	-0.576	0.565	-17.830	9.736
<b>x2_Bushwick</b>	-2.7252	7.032	-0.388	0.698	-16.507	11.057
<b>x2_Chelsea</b>	-3.4508	7.122	-0.484	0.628	-17.411	10.509
<b>x2_Crown Heights</b>	-4.8392	7.061	-0.685	0.493	-18.678	9.000
<b>x2_East Flatbush</b>	-2.0758	7.246	-0.286	0.775	-16.277	12.125
<b>x2_East Harlem</b>	-3.6134	7.118	-0.508	0.612	-17.564	10.337
<b>x2_East Village</b>	-5.0527	7.046	-0.717	0.473	-18.862	8.757
<b>x2_Financial District</b>	-2.9938	7.150	-0.419	0.675	-17.009	11.021
<b>x2_Flatbush</b>	-2.7533	7.187	-0.383	0.702	-16.839	11.332
<b>x2_Greenpoint</b>	-4.1062	7.098	-0.578	0.563	-18.019	9.806
<b>x2_Harlem</b>	-2.3843	7.064	-0.338	0.736	-16.230	11.461
<b>x2_Hell's Kitchen</b>	-4.1698	7.089	-0.588	0.556	-18.064	9.724
<b>x2_Lower East Side</b>	-3.8904	7.133	-0.545	0.586	-17.872	10.091
<b>x2_Midtown</b>	-4.9726	7.119	-0.699	0.485	-18.925	8.980
<b>x2_Other</b>	-4.1062	7.000	-0.587	0.557	-17.826	9.613
<b>x2_Prospect-Lefferts Gardens</b>	-5.3995	7.206	-0.749	0.454	-19.523	8.724
<b>x2_Upper East Side</b>	-6.3507	7.075	-0.898	0.369	-20.217	7.515
<b>x2_Upper West Side</b>	-4.6743	7.049	-0.663	0.507	-18.491	9.143
<b>x2_Washington Heights</b>	-2.5956	7.136	-0.364	0.716	-16.582	11.391

<b>x2_West Village</b>	-3.2539	7.171	-0.454	0.650	-17.308	10.800
<b>x2_Williamsburg</b>	-4.7611	7.002	-0.680	0.497	-18.485	8.963
<b>service fee</b>	4.9429	0.003	1706.296	0.000	4.937	4.949
<b>minimum nights</b>	0.0083	0.006	1.287	0.198	-0.004	0.021
<b>number of reviews</b>	-0.0083	0.005	-1.755	0.079	-0.017	0.001
<b>reviews per month</b>	0.2119	0.138	1.533	0.125	-0.059	0.483
<b>review rate number</b>	-0.0007	0.151	-0.004	0.996	-0.298	0.296
<b>calculated host listings count</b>	-0.0120	0.007	-1.635	0.102	-0.026	0.002
<b>availability 365</b>	0.0003	0.001	0.198	0.843	-0.003	0.003
<b>lr_year</b>	0.0800	0.133	0.603	0.547	-0.180	0.340
<b>lr_month</b>	-0.0983	0.069	-1.426	0.154	-0.233	0.037
<b>lr_day</b>	0.0553	0.020	2.798	0.005	0.017	0.094

<b>Omnibus:</b>	65755.042	<b>Durbin-Watson:</b>	1.997
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	343304984.137
<b>Skew:</b>	3.674	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	360.300	<b>Cond. No.</b>	1.00e+16

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 2.65e-21. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

From these regression results, it looks like variables that effect the price the most are unconfirmed, verified, or choosing a location part of the 5 boroughs. There is an average of about 89 dollars per AirBnB. Some neighborhoods that would be cheaper than others would be Upper East Side, Astoria, Prospect-Lefferts Gardens, and East Village. It seems as though time of the year to get an AirBnB does not affect the price that much. The location you choose seems to affect it more.

## ▼ 4.5 Building Regressor Tree Models

Multiple models will be created to see which one has the lowest RMSE.

### ▼ 4.5.1 Baseline Decision Tree

Need to start off with a baseline before using more complex models.


```
In [90]: # Instantiating and fitting a DecisionTreeRegressor
# Get predictions for training
baseline_tree = DecisionTreeRegressor(random_state=123)
baseline_tree.fit(data_train, target_train)
predictions = baseline_tree.predict(data_test)
```

```
In [91]: # Examining how important each feature ends up being
# Feature importance
baseline_tree.feature_importances_
```

```
Out[91]: array([1.49562469e-04, 2.14342685e-04, 1.45394485e-06, 2.67307389e-04,
1.37657987e-04, 1.34919410e-04, 1.27152516e-05, 1.52206610e-05,
2.50950456e-04, 1.19473540e-04, 4.20819088e-05, 1.11187182e-05,
4.13553821e-08, 9.07184853e-05, 1.46107985e-04, 2.11114211e-07,
6.88450508e-06, 1.37050303e-05, 6.18981071e-06, 8.82118729e-05,
1.31773800e-05, 1.79347159e-04, 1.04258574e-04, 3.04305914e-06,
8.37332607e-05, 4.12346335e-05, 2.03533480e-05, 6.32259640e-07,
1.44733812e-04, 9.88443880e-01, 8.94427513e-04, 1.41545222e-03,
1.59956555e-03, 7.55372163e-04, 6.29478049e-04, 1.25991650e-03,
2.56711372e-04, 5.81344304e-04, 1.86446430e-03])
```

```
In [92]: # Visualizing feature importance
def plot_feature_importances(model):
    n_features = data_train.shape[1]
    plt.figure(figsize=(8,8))
    plt.barh(range(n_features), model.feature_importances_, align='center')
    plt.yticks(np.arange(n_features), data_train.columns.values)
    plt.xlabel('Feature importance')
    plt.ylabel('Feature')

plot_feature_importances(baseline_tree)
```



Service fee seems to be the only important feature.

```
In [93]: # Test set predictions
baseline_pred = baseline_tree.predict(data_test)
```

```
In [94]: # Evaluating the Algorithm
# actual vs predicted
print('Mean Absolute Error:', metrics.mean_absolute_error(target_test, baseline_pred))
print('Mean Squared Error:', metrics.mean_squared_error(target_test, baseline_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(target_test, baseline_p

Mean Absolute Error: 5.392583120204604
Mean Squared Error: 3335.096361311323
Root Mean Squared Error: 57.750293170782456
```

## ▼ 4.5.2 Random Forest

```
In [95]: # Instantiate and fit RandomForestRegressor
regressor = RandomForestRegressor(n_estimators = 50, random_state = 123)
```

Passing training data to train the RF regressor model

```
In [96]: regressor.fit(data_train, target_train)
```

```
Out[96]: RandomForestRegressor(n_estimators=50, random_state=123)
```

Predicting test set results using RF regressor model


```
In [97]: # Predicting the Test set results
target_pred = regressor.predict(data_test)
```

Evaluate RF regressor model algorithm using error metrics.

```
In [98]: # Evaluating the Algorithm
# With n_estimators = 50
# actual vs predicted
print('Mean Absolute Error:', metrics.mean_absolute_error(target_test, target_pred))
print('Mean Squared Error:', metrics.mean_squared_error(target_test, target_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(target_test, target_pre
```

```
Mean Absolute Error: 5.099519755095713
Mean Squared Error: 1821.9673312665452
Root Mean Squared Error: 42.6845092658513
```

```
In [99]: plot_feature_importances(regressor)
```



Trying n\_estimator = 40

```
In [100]: regressor2 = RandomForestRegressor(n_estimators = 40, random_state = 123)
```

```
In [101]: regressor2.fit(data_train, target_train)
# Predicting the Test set results
target_pred2 = regressor2.predict(data_test)
```

```
In [102]: # Evaluating the Algorithm
# With n_estimators = 40
# actual vs predicted
print('Mean Absolute Error:', metrics.mean_absolute_error(target_test, target_pred2))
print('Mean Squared Error:', metrics.mean_squared_error(target_test, target_pred2))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(target_test, target_pre
```

```
Mean Absolute Error: 5.129389889948073
Mean Squared Error: 1845.9050140644972
Root Mean Squared Error: 42.96399671893314
```

Trying n\_estimator = 30

```
In [103]: regressor3 = RandomForestRegressor(n_estimators = 30, random_state = 123)
```

```
In [104]: regressor3.fit(data_train, target_train)
# Predicting the Test set results
target_pred3 = regressor3.predict(data_test)
```

```
In [105]: # Evaluating the Algorithm
# With n_estimators = 30
# actual vs predicted
print('Mean Absolute Error:', metrics.mean_absolute_error(target_test, target_pred3))
print('Mean Squared Error:', metrics.mean_squared_error(target_test, target_pred3))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(target_test, target_pre
```

```
Mean Absolute Error: 5.146629513992686
Mean Squared Error: 1864.4637599824623
Root Mean Squared Error: 43.17943677240895
```

### ▼ 4.5.3 Random Forest with Grid Search

```
In [106]: # Number of trees in random forest
# Create the parameter grid based on the results of random search
param_grid = {
    'bootstrap': [True],
    'max_depth': [80, 90],
    'max_features': [2, 3],
    'min_samples_leaf': [3, 4],
    'min_samples_split': [10, 12],
    'n_estimators': [200, 300]
}
# Create a based model
rf_grid = RandomForestRegressor()
# Instantiate the grid search model
grid_search = GridSearchCV(estimator = rf_grid, param_grid = param_grid,
                           cv = 3, n_jobs = -1, verbose = 2)
```

```
In [107]: # Fit grid search to data
grid_search.fit(data_train, target_train)
grid_search.best_params_
{'bootstrap': True,
 'max_depth': 80,
 'max_features': 3,
 'min_samples_leaf': 5,
 'min_samples_split': 10,
 'n_estimators': 300}
best_grid = grid_search.best_estimator_
```

Fitting 3 folds for each of 32 candidates, totalling 96 fits

[Parallel(n\_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

[Parallel(n\_jobs=-1)]: Done 25 tasks | elapsed: 55.2s

[Parallel(n\_jobs=-1)]: Done 96 out of 96 | elapsed: 3.2min finished

```
In [108]: print("The best hyperparameters are ", grid_search.best_params_)
```

The best hyperparameters are {'bootstrap': True, 'max\_depth': 80, 'max\_features': 3, 'min\_samples\_leaf': 3, 'min\_samples\_split': 10, 'n\_estimators': 200}

Using best hyperparameters for final calculation.

```
In [109]: rf_grid = RandomForestRegressor(bootstrap = grid_search.best_params_["bootstrap"],
                                         max_depth = grid_search.best_params_["max_depth"],
                                         max_features = grid_search.best_params_["max_features"],
                                         min_samples_leaf = grid_search.best_params_["min_samples_leaf"],
                                         min_samples_split = grid_search.best_params_["min_samples_split"],
                                         n_estimators = grid_search.best_params_["n_estimators"])
rf_grid.fit(data_train, target_train)
```

```
Out[109]: RandomForestRegressor(max_depth=80, max_features=3, min_samples_leaf=3,
                               min_samples_split=10, n_estimators=200)
```

```
In [110]: rf_pred = rf_grid.predict(data_test)
```

```
In [111]: # Evaluating the Algorithm
# actual vs predicted
print('Mean Absolute Error:', metrics.mean_absolute_error(target_test, rf_pred))
print('Mean Squared Error:', metrics.mean_squared_error(target_test, rf_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(target_test, rf_pred)))
```


Mean Absolute Error: 75.46257790949845

Mean Squared Error: 10230.180315716561

Root Mean Squared Error: 101.14435384991374



```
In [112]: plot_feature_importances(rf_grid)
```



Similar the other models above, service fee still seems way higher than the rest.

#### ▼ 4.5.4 XG Boost

```
In [113]: import xgboost as xgb
```

```
In [114]: # Instantiate XGBClassifier
xg_regressor = xgb.XGBRegressor(eval_metric='rmse')
```

```
In [127]: # Creating second one without service fee
xg_regressor2 = xgb.XGBRegressor(eval_metric='rmse')
```

```
In [334]: # Creating third one with just main feaatures
xg_regressor3 = xgb.XGBRegressor(eval_metric='rmse')
```

```
In [115]: # Set up search grid
param_grid = {"max_depth": [4, 5],
              "n_estimators": [100, 200, 300],
              "learning_rate": [0.01, 0.015]}
```

```
In [116]: # Trying out every combination of the above values
search = GridSearchCV(xg_regressor, param_grid, cv=5).fit(data_train, target_train)
print("The best hyperparameters are ", search.best_params_)
```

The best hyperparameters are {'learning\_rate': 0.015, 'max\_depth': 4, 'n\_estimators': 300}

Using best hyperparameters for final calculation.

```
In [117]: xg_regressor = xgb.XGBRegressor(learning_rate = search.best_params_["learning_rate"],
                                          n_estimators = search.best_params_["n_estimators"],
                                          max_depth = search.best_params_["max_depth"],)

xg_regressor.fit(data_train, target_train)
```

```
Out[117]: XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                        colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
                        importance_type='gain', interaction_constraints='',
                        learning_rate=0.015, max_delta_step=0, max_depth=4,
                        min_child_weight=1, missing=nan, monotone_constraints='()',
                        n_estimators=300, n_jobs=0, num_parallel_tree=1, random_state=0,
                        reg_alpha=0, reg_lambda=1, scale_pos_weight=1, subsample=1,
                        tree_method='exact', validate_parameters=1, verbosity=None)
```


Using model for predictions.

```
In [118]: xg_pred = xg_regressor.predict(data_test)
```

```
In [119]: # Evaluating the Algorithm
# actual vs predicted
print('Mean Absolute Error:', metrics.mean_absolute_error(target_test, xg_pred))
print('Mean Squared Error:', metrics.mean_squared_error(target_test, xg_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(target_test, xg_pred)))
```

Mean Absolute Error: 10.570509811823547  
Mean Squared Error: 1577.1709730373932  
Root Mean Squared Error: 39.713611936430475

```
In [120]: plot_feature_importances(xg_regressor)
```



Out of all the models used, XGBoost had the lowest RMSE which was 39.71

#### ▼ XG Bosst without Service Fee

```
In [128]: # Trying out every combination of param grid
search2 = GridSearchCV(xg_regressor2, param_grid, cv=5).fit(data_train2, target_train2)
print("The best hyperparameters are ", search2.best_params_)
```

The best hyperparameters are {'learning\_rate': 0.015, 'max\_depth': 4, 'n\_estimators': 300}

```
In [129]: xg_regressor2 = xgb.XGBRegressor(learning_rate = search2.best_params_["learning_rate"],
                                           n_estimators = search2.best_params_["n_estimators"],
                                           max_depth = search2.best_params_["max_depth"],)

xg_regressor2.fit(data_train2, target_train2)
```

```
Out[129]: XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                        colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
                        importance_type='gain', interaction_constraints='',
                        learning_rate=0.015, max_delta_step=0, max_depth=5,
                        min_child_weight=1, missing=nan, monotone_constraints='()',
                        n_estimators=300, n_jobs=0, num_parallel_tree=1, random_state=0,
                        reg_alpha=0, reg_lambda=1, scale_pos_weight=1, subsample=1,
                        tree_method='exact', validate_parameters=1, verbosity=None)
```

```
In [130]: xg_pred2 = xg_regressor2.predict(data_test2)
```


```
In [131]: # Evaluating the Algorithm
# actual vs predicted
print('Mean Absolute Error:', metrics.mean_absolute_error(target_test2, xg_pred2))
print('Mean Squared Error:', metrics.mean_squared_error(target_test2, xg_pred2))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(target_test2, xg_pred2)))
```

```
Mean Absolute Error: 287.55282549207305
Mean Squared Error: 110496.0780582735
Root Mean Squared Error: 332.40950356190706
```

RMSE increased exponentially from ~39 to ~332.

```
In [133]: # Rewriting plot feature importance function
# Visualizing feature importance
def plot_feature_importances2(model):
    n_features = data_train2.shape[1]
    plt.figure(figsize=(8,8))
    plt.barh(range(n_features), model.feature_importances_, align='center')
    plt.yticks(np.arange(n_features), data_train2.columns.values)
    plt.xlabel('Feature importance')
    plt.ylabel('Feature')
```

```
In [134]: plot_feature_importances2(xg_regressor2)
```



Featured importance data is all over the place. Top 6 features according to this chart are West Village, Bedford-Stuyvesant, Upper West Side, East Village, East Flatbush, Prospect-Lefferts Gardens. Among the continuous data the top features seem to be availability\_365, lr\_year, reviews per month, and calculated host listings count

#### ▼ XG Boost with just main features

```
In [335]: # Trying out every combination of param grid
search3 = GridSearchCV(xg_regressor3, param_grid, cv=5).fit(data_train3, target_train3)
print("The best hyperparameters are ", search3.best_params_)
```

The best hyperparameters are {'learning\_rate': 0.015, 'max\_depth': 5, 'n\_estimators': 300}

```
In [336]: xg_regressor3 = xgb.XGBRegressor(learning_rate = search3.best_params_["learning_rate"],
                                           n_estimators = search3.best_params_["n_estimators"],
                                           max_depth = search3.best_params_["max_depth"],)

xg_regressor3.fit(data_train3, target_train3)
```

```
Out[336]: XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                        colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
                        importance_type='gain', interaction_constraints='',
                        learning_rate=0.015, max_delta_step=0, max_depth=5,
                        min_child_weight=1, missing=nan, monotone_constraints=(),
                        n_estimators=300, n_jobs=0, num_parallel_tree=1, random_state=0,
                        reg_alpha=0, reg_lambda=1, scale_pos_weight=1, subsample=1,
                        tree_method='exact', validate_parameters=1, verbosity=None)
```


```
In [337]: xg_pred3 = xg_regressor3.predict(data_test3)
```

```
In [338]: # Evaluating the Algorithm
# actual vs predicted
print('Mean Absolute Error:', metrics.mean_absolute_error(target_test3, xg_pred3))
print('Mean Squared Error:', metrics.mean_squared_error(target_test3, xg_pred3))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(target_test3, xg_pred3)))
```

Mean Absolute Error: 287.68327523475193  
Mean Squared Error: 110623.00374710109  
Root Mean Squared Error: 332.6003664265887

```
In [342]: # Rewriting plot feature importance function
# Visualizing feature importance
def plot_feature_importances3(model):
    n_features = data_train3.shape[1]
    plt.figure(figsize=(8,8))
    plt.barh(range(n_features), model.feature_importances_, align='center')
    plt.yticks(np.arange(n_features), data_train3.columns.values)
    plt.xlabel('Feature importance')
    plt.ylabel('Feature')
```

```
In [343]: plot_feature_importances3(xg_regressor3)
```



From this reduced feature importance graph, it shows that an Airbnb being available for 365 days a year is most important to the price. Bookings in Queens seem to also impact the price the most out of all the boroughs. It also seems that booking on a certain day impacts price more slightly than booking on a certain month or year. Reviews per month of a Airbnb equally impacts the price as well. Finally, being an unconfirmed host seems to have more of an impact on the price than a verified host.

## 5 Conclusion

After analyzing this dataset from Airbnb, I got a clearer picture of what customers prefer and what effects the total price. To start off with, here is some general information. The most stayed at neighborhood seemed to be Bedford-Stuyvesant (Bed-Stuy) in Brooklyn. Secondly, some of the most expensive neighborhoods people stay at on average are Flatbush,

West Village, East Flatbush, Prospect-Lefferts Gardens, and Astoria. People tend to stay their longest visits in Upper East Side (11.5) days, Upper West Side, Midtown, Chelsea, and Hell's Kitchen. Building off that, the data suggests that people spend the longest stays in Manhattan on average compared to the rest of New York's boroughs. Speaking of booking, the amount of nights people tend to stay at a location is 1-2 nights. Diving further into the data, a OLS Regression chart was created. From this, it was seen that the average price for an Airbnb amongst customers was \$89. The neighborhoods that someone could stay at for a cheaper price are Upper East Side, Astoria, Prospect-Lefferts Gardens, and East Village. Finally, it seems that the service fee for these stays predictably affect the total price the most for an Airbnb. Finally the data was applied to several regressor tree models. Those being a baseline one, random forest with different number of estimators, random forest with a grid search applied, and a XG Boost model with Grid Search. Among these variations, the model with the lowest RMSE was XG Boost with a value of 39.71. From all this data collected, one can say people tend to use Airbnbs to stay in Manhattan and their visits are longer than anywhere else. If one wanted to stay at the most affordable areas, they would want to stay in either Upper East Side, Astoria, Prospect-Lefferts Gardens, and East Village. Finally, the host of these Airbnbs should expect the average person to stay 1-2 nights.